

Zachary Robertson

GRADUATE STUDENT

✉ zroberts@stanford.edu | 🌐 zrobertson466920.github.io/ | Phone: (832)-318-3039 | Mail: 51 Dudley Ln., Apt. 329 Stanford, CA 94305

"Play is the highest form of research."

Education

Stanford

PH.D. IN COMPUTER SCIENCE

- Google Graduate Fellowship in Computer Science
- EDGE Fellowship
- NSF Honorary Mention

Stanford, U.S.

Sept. 2022 -

University of Illinois at Urbana-Champaign

MS IN COMPUTER SCIENCE

- Thesis on Performance Metric Elicitation
- GEM fellowship
- Wing Kai Cheng Fellowship

Champaign, U.S.

Aug. 2020 - May. 2022

University of Chicago

B.S. IN COMPUTATIONAL AND APPLIED MATH

- Graduated with Honors
- Jackie Robinson Scholarship

Chicago, U.S.

Oct. 2016 - June 2020

LoneStar Community College

ASSOCIATE DEGREE IN SCIENCE

- President's List
- Home schooled with dual enrollment at local community college

Spring, TX, U.S.

Oct. 2013 - May. 2016

Experience

Stanford University

INSTRUCTOR

- Instructor for Introduction to AI Safety (CS120); previously taught core Artificial Intelligence curriculum (CS221)

Stanford, U.S.

2025 - Present

STAIR Lab at Stanford (advised by Sanmi Koyejo)

RESEARCH ASSISTANT

- Developed mutual evaluation, an information-theoretic framework for evaluating AI systems without ground truth; first-author work published in TMLR
- Conceived, wrote, and submitted a \$500K OpenAI Superalignment grant proposal; advisor served as institutional PI

Stanford, U.S.

Sept. 2022 - Present

Google

STUDENT RESEARCHER

- Designed tractable surrogates for welfare maximization
- Applied these methods to predicted click-through-rate models used in Ads

Mountain View, U.S.

Apr. 2022 - Aug. 2022

Lam Research

RESEARCH INTERN

- Applied reinforcement learning and machine learning to optimize semiconductor wafer production flows
- Used Python, Open AI Gym, and Tensorflow

Fremont, U.S.

June 2020 - March 2021

Robot Intelligence through Perception Lab at TTIC

RESEARCH ASSISTANT

- Implemented a model to do sparse-depth completion on a robot using PyTorch, Docker, and SLURM.
- Trained a simulated UR5 arm to pick up blocks from a table with natural language using PyTorch, OpenGym, and Mujoco.
- Engaged in theoretical study of reinforcement learning using regret based approach for Honor's Thesis

Chicago, U.S.

Jun. 2019 - June 2020

Allen and Company LLC.

ANALYST

- Evaluated quantum-computing startups and cryptocurrencies, automated internal reporting, and analyzed investor sentiment

New York, U.S.

June 2018 - Aug. 2018

Honors & Awards

2024	Idea and Writing , Lead author, Superalignment Fast Grant proposal (\$500k, OpenAI) awarded to STAIR lab	Stanford
2023	Idea and Writing , Accelerating Foundation Models Research (Microsoft) awarded to STAIR lab	Stanford
2022	Recipient , Google Graduate Fellowship in Computer Science	Stanford
2022	Recipient , EDGE Fellowship	Stanford
2021	Honorary Mention , NSF GRFP Fellowship	UIUC
2021	Recipient , Wing Kai Cheng Fellowship	UIUC
2020	Recipient , GEM Fellowship	UChicago
2018	Recipient , Dean's List	UChicago
2017	Winner , Illinois Blockchain Hackathon	Chicago/Urbana
2016	Scholar , Jackie Robinson Foundation	New York

Service

WellLabeled Affinity Group

LEAD

HAI
Sept. 2023 - May. 2024

- Led an advocacy group for data annotation worker rights
- Interviewed stakeholders from Turkopticon, Scale AI, and OpenAI
- Featured article and video testimonial. [See here](#).

Alignment Newsletter (Organized by Rohin Shah)

CONTRIBUTOR

CHAI
Sept. 2019 - Sept. 2023

- Wrote summaries and opinions on advances in machine learning for the Alignment Newsletter
- Over 1k subscribers

Skills

Programming	Python, LaTeX, C++, Haskell
Machine Learning	Evaluation design, mechanism design, adversarial robustness
Math	Analysis, Linear Algebra, Statistical Learning Theory, Optimization, Statistics

Publications and Presentations

Presentations and External Impact

1. "Paying for Information: Incentive-Compatible LLM Evaluation via f-Mutual Information", [Invited Lightning Talk](#), *Building an Aligned AI Future* (Fifty Years @ Stanford), September 2025
2. "Information-Theoretic Evaluation Methods", [Featured Implementation](#), Jina AI Production Systems (\$30M Series A), 2024
3. "Exploring the Complex Ethical Challenges of Data Annotation", [Invited Talk](#), FloodGate, San Francisco, September 2024
4. "Towards Scalable Information Elicitation for Oversight in Human-AI Systems", Invited Talk, Max Planck Institute for Intelligent Systems, Tübingen, July 2024

Selected Publications

1. Farnaz Jahanbakhsh, Dora Zhao, Tiziano Piccardi, Zachary Robertson, Ziv Epstein, Sanmi Koyejo, and Michael S Bernstein. Value alignment of social media ranking algorithms. *arXiv preprint arXiv:2509.14434*, 2025
2. Zachary Robertson and Sanmi Koyejo. Let's measure information step-by-step: Llm-based evaluation beyond vibes. *arXiv preprint arXiv:2508.05469*, 2025. Accepted with minor revisions at Transactions on Machine Learning Research (TMLR), 2026
3. Zachary Robertson and Sanmi Koyejo. Implicit regularization in feedback alignment mechanisms. *ICML*, 2024
4. Boxiang Lyu, Zhe Feng, Zachary Robertson, and Sanmi Koyejo. Pairwise ranking losses of click-through rates prediction for welfare maximization in ad auctions. *ICML*, 2023
5. Zachary Robertson, Hantao Zhang, and Sanmi Koyejo. Cooperative inverse decision theory for uncertain preferences. *AISTATS*, 2023